

```
import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from torch.utils.data import TensorDataset, DataLoader
```

```
df = pd.read_csv("/content/customers.csv")
print("Dataset Shape:", df.shape)
print(df.head())
```

Dataset Shape: (10695, 11)

	ID	Gender	Ever_Married	Age	Graduated	Profession	Work_Experience	\
0	462809	Male	No	22	No	Healthcare	1.0	
1	462643	Female	Yes	38	Yes	Engineer	NaN	
2	466315	Female	Yes	67	Yes	Engineer	1.0	
3	461735	Male	Yes	67	Yes	Lawyer	0.0	
4	462669	Female	Yes	40	Yes	Entertainment	NaN	

	Spending_Score	Family_Size	Var_1	Segmentation
0	Low	4.0	Cat_4	D
1	Average	3.0	Cat_4	A
2	Low	1.0	Cat_6	B
3	High	2.0	Cat_6	B
4	High	6.0	Cat_6	A

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10695 entries, 0 to 10694
Data columns (total 11 columns):
#   Column              Non-Null Count  Dtype
---  -
0   ID                   10695 non-null  int64
1   Gender               10695 non-null  object
2   Ever_Married         10505 non-null  object
3   Age                  10695 non-null  int64
4   Graduated            10593 non-null  object
5   Profession            10533 non-null  object
6   Work_Experience       9597 non-null   float64
7   Spending_Score       10695 non-null  object
8   Family_Size          10247 non-null  float64
9   Var_1                10587 non-null  object
10  Segmentation         10695 non-null  object
dtypes: float64(2), int64(2), object(7)
memory usage: 919.2+ KB
```

```
df.describe()
```

	ID	Age	Work_Experience	Family_Size
<b>count</b>	10695.000000	10695.000000	9597.000000	10247.000000
<b>mean</b>	463468.088640	43.511828	2.619777	2.844052
<b>std</b>	2600.966411	16.774158	3.390790	1.536427
<b>min</b>	458982.000000	18.000000	0.000000	1.000000
<b>25%</b>	461220.500000	30.000000	0.000000	2.000000
<b>50%</b>	463451.000000	41.000000	1.000000	3.000000
<b>75%</b>	465733.500000	53.000000	4.000000	4.000000
<b>max</b>	467974.000000	89.000000	14.000000	9.000000

```
df = df.drop(columns=["ID"])
```

```
df.fillna({"Work_Experience": 0, "Family_Size": df["Family_Size"].median()}, inplace=True)
```

```
categorical_columns = ["Gender", "Ever_Married", "Graduated", "Profession", "Spending_Score", "Var_1"]
for col in categorical_columns:
    df[col] = LabelEncoder().fit_transform(df[col])
```

```
label_encoder = LabelEncoder()
df["Segmentation"] = label_encoder.fit_transform(df["Segmentation"])
```

```

X = df.drop(columns=["Segmentation"])
y = df["Segmentation"].values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

```

```

X_train = torch.tensor(X_train, dtype=torch.float32)
X_test = torch.tensor(X_test, dtype=torch.float32)
y_train = torch.tensor(y_train, dtype=torch.long)
y_test = torch.tensor(y_test, dtype=torch.long)

train_dataset = TensorDataset(X_train, y_train)
test_dataset = TensorDataset(X_test, y_test)
train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=16)

```

```

class PeopleClassifier(nn.Module):
    def __init__(self, input_size):
        super(PeopleClassifier, self).__init__()
        self.fc1 = nn.Linear(input_size, 32)
        self.fc2 = nn.Linear(32, 20)
        self.fc3 = nn.Linear(20, 16)
        self.fc4 = nn.Linear(16, 4)

    def forward(self, x):
        x=F.relu(self.fc1(x))
        x=F.relu(self.fc2(x))
        x=F.relu(self.fc3(x))
        x=self.fc4(x)
        return x

```

```

def train_model(model,train_loader,criterion,optimizer,epochs=100):
    for epoch in range(epochs):
        model.train()
        for X_batch,y_batch in train_loader:
            optimizer.zero_grad()
            outputs=model(X_batch)
            loss=criterion(outputs,y_batch)
            loss.backward()
            optimizer.step()

        if(epoch+1)%10==0:
            print(f'Epoch [{epoch+1}/{epochs}],Loss:{loss.item():.4f}')

```

```

model =PeopleClassifier(input_size=X_train.shape[1])
criterion =nn.CrossEntropyLoss()
optimizer =optim.Adam(model.parameters(),lr=0.001)

train_model(model,train_loader,criterion,optimizer,epochs=100)

```

```
Epoch [100/100],Loss:0.7839
```

```

model.eval()
predictions, actuals = [], []
with torch.no_grad():
    for X_batch, y_batch in test_loader:
        outputs = model(X_batch)
        _, predicted = torch.max(outputs, 1)
        predictions.extend(predicted.numpy())
        actuals.extend(y_batch.numpy())

```

```

accuracy = accuracy_score(actuals, predictions)
conf_matrix = confusion_matrix(actuals, predictions)
class_report = classification_report(actuals, predictions, target_names=[str(i) for i in label_encoder.classes_])
print("Name: A.LAHARI")
print("Register No: 212223230111")
print(f'Test Accuracy: {accuracy:.2f}%')
print("Confusion Matrix:\n", conf_matrix)
print("Classification Report:\n", class_report)

```

```

Name: A.LAHARI
Register No: 212223230111
Test Accuracy: 0.46%
Confusion Matrix:
[[192 118  98 176]

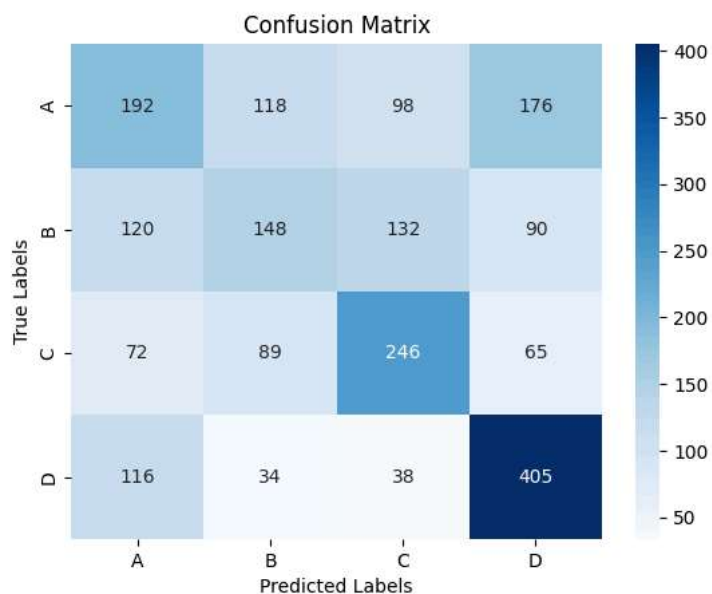
```

```
[120 148 132 90]
[ 72  89 246 65]
[116  34  38 405]]
Classification Report:
              precision    recall  f1-score   support

     A       0.38       0.33       0.35         584
     B       0.38       0.30       0.34         490
     C       0.48       0.52       0.50         472
     D       0.55       0.68       0.61         593

 accuracy          0.46         2139
 macro avg         0.45         2139
 weighted avg      0.45         2139
```

```
import seaborn as sns
import matplotlib.pyplot as plt
sns.heatmap(conf_matrix, annot=True, cmap='Blues', xticklabels=label_encoder.classes_, yticklabels=label_encoder.classes_, fmt:
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.title("Confusion Matrix")
plt.show()
```



```
sample_input = X_test[9].clone().unsqueeze(0).detach().type(torch.float32)
with torch.no_grad():
    output = model(sample_input)
    predicted_class_index = torch.argmax(output[0]).item()
    predicted_class_label = label_encoder.inverse_transform([predicted_class_index])[0]
print("Name: A.LAHARI")
print("Register No: 21223230111")
print(f'Predicted class for sample input: {predicted_class_label}')
print(f'Actual class for sample input: {label_encoder.inverse_transform([y_test[9].item()])[0]}')
```

```
Name: A.LAHARI
Register No: 21223230111
Predicted class for sample input: D
Actual class for sample input: A
```

